


Views



When you want to provide limited access to the existing data, then we create views

views are of 2 types

1. views

- for view no separate memory is allocated for storing data,
- only base query will be stored, and any statement on the view will use base query to get the data, because of that we always get UpToDate data in view
- if view contains all not null columns of the single base table and if it is not read only view, then we can use all DML operation(insert, delete, update) on the view
- if the view is based on joins, aggregate functions, group by statement or union of multiple queries, then by default the views are readonly

```
select * from emp e
where sal < (select avg(sal) from emp m where m.mgr=e.mgr)
```

```
create a view fac_room
as
select cid,cname,fid,fname,null,null
from course c right join faculty f on f.fid=c.fid
where c.cname is null
union
select cid,cname,null,null,rid,rname
from course c right join room r on r.rid=c.rid
where c.cname is null
```

uses of views:

- Hide complexity of the queries (joins, aggregate functions nested queries)
- To give restricted access to few columns or rows from tables
- Hide table names, to increase the security of data.

2. Materialized view

- a. Views for which the first time the base query will get executed and then the output will be stored in a temporary table in RAM, within the session, the data will be retrieved from the RAM
- b. you may not get upto date data in materialized view.

create materialized view myview

as

select * from emp;

to drop view

drop view myview

to get the 3rd highest salary

select *

-> from emp s

-> where 2=(select count(*)

-> from (select distinct sal from emp) e

-> where e.sal>s.sal)

-> order by sal;

to get 3 topmost salaried employee

select *

-> from emp s

-> where 3>select count(*)

-> from (select distinct sal from emp) e

-> where e.sal>s.sal)

-> order by sal;

to get 3 bottommost salaried employee

select *

-> from emp s

-> where 3>select count(*)

-> from (select distinct sal from emp) e

-> where e.sal<s.sal)

-> order by sal;

to create index

```
create index sal_idx  
on emp(sal)
```

to drop index

```
drop index sal_idx
```

to list all indexes available on the table

```
SHOW INDEXES FROM customers
```

to find which indexes are used by the query or to suggest using of which indexes in query

```
EXPLAIN SELECT *
```

```
FROM
```

```
customers
```

```
WHERE
```

```
contactFirstName LIKE 'A%'
```

```
OR contactLastName LIKE 'A%';
```

MySQL Engine INNODB supports ACID property

A—atomicity-> every transaction will get executed as a single unit

```
begin transaction
```

```
    withdraw amt form source
```

```
    deposit amt in the ddestination
```

```
end transaction
```

C—consistency---correctness of data will be there after every transaction completes

I ----isolation---- every transaction will get executed in isolation,

intermediate changes are visible only to the user who is performing the transaction,

these changes will visible to other users, when the transaction is committed

D- durability-----It will show the performance for longer period of time

Transaction control language

commit, rollback, savepoint

emp

10 records

commit

insert 1

insert 1

update 1

commit

delete 1

rollback

10 records

commit

insert-3

update 2

savepoint A

delete 1

insert-3

update 2

savepoint B

insert-2

delete 1

rollback to B

12 rows

insert -3

delete 1

create table

insert- 4

delete 1

rollback

